

- [Product Service Deployment Guide](#)
 - [Deployment Options](#)
 - [Local Deployment](#)
 - [Prerequisites](#)
 - [Steps](#)
 - [Docker Deployment](#)
 - [Prerequisites](#)
 - [Steps](#)
 - [Docker Compose Components](#)
 - [Docker Network](#)
 - [Volume Management](#)
 - [CI/CD Pipeline Integration](#)
 - [Container Optimization Techniques](#)
 - [Cloud Deployment Considerations](#)
 - [AWS Deployment](#)
 - [Azure Deployment](#)
 - [Google Cloud Deployment](#)
 - [Scaling Considerations](#)

Product Service Deployment Guide

Deployment Options

The Product Service supports multiple deployment options to accommodate various infrastructure requirements:

1. **Local Development** - Run directly on developer machine
2. **Containerized Development** - Using Docker and Docker Compose
3. **Kubernetes Deployment** - Production-ready orchestration (configuration provided)
4. **Cloud Platform Deployment** - AWS ECS and Azure AKS templates available

Local Deployment

Prerequisites

- Java 17 JDK
- Maven 3.8+
- PostgreSQL 13+ (optional, can use H2 for development)
- Kafka (optional, can be disabled for development)

Steps

1. Clone the repository
2. Configure the application:

```
# Edit application.yml with your local settings
```

3. Build the application:

```
mvn clean package
```

4. Run the application:

```
java -jar target/product-service.jar
```

Docker Deployment

Prerequisites

- Docker 20.10+
- Docker Compose 2.0+

Steps

1. Navigate to the project directory:

```
cd ecommerce-microservices/product-service
```

2. Build and run with Docker Compose:

```
docker-compose up --build
```

3. Access the services:

- Product Service API: <http://localhost:8080/api>
- Kafka UI: <http://localhost:8090>
- PostgreSQL: <localhost:5432>

Docker Compose Components

The `docker-compose.yml` file orchestrates the following services:

1. Zookeeper:

- Image: [confluentinc/cp-zookeeper:7.3.0](#)
- Purpose: Manages Kafka cluster state
- Port: 2181

2. Kafka:

- Image: [confluentinc/cp-kafka:7.3.0](#)
- Purpose: Event streaming platform
- Ports: 9092, 29092
- Dependencies: Zookeeper

3. Kafka UI:

- Image: [provectuslabs/kafka-ui:latest](#)
- Purpose: Web interface for Kafka monitoring
- Port: 8090
- Dependencies: Kafka

4. PostgreSQL:

- Image: [postgres:13](#)

- Purpose: Relational database
- Port: 5432
- Volumes: `postgres_data`

5. Product Service:

- Image: Built from Dockerfile
- Purpose: Spring Boot microservice
- Port: 8080
- Dependencies: PostgreSQL, Kafka

Docker Network

All services are connected to a custom bridge network called `product-network`, allowing containers to communicate using service names as hostnames.

Volume Management

The PostgreSQL data is persisted using a named volume `postgres_data`, ensuring data survival across container restarts.

CI/CD Pipeline Integration

The service includes GitHub Actions workflow configurations for automated CI/CD:

1. Build and Test:

- Triggers on pull requests to main branch
- Builds the application and runs unit/integration tests
- Performs static code analysis with SonarQube

2. Deploy to Development:

- Triggers on merges to main branch
- Builds Docker image
- Pushes to container registry
- Deploys to development environment

3. Deploy to Production:

- Triggers on release tags
- Builds production Docker image
- Pushes to production container registry
- Deploys to production environment with approval step

Container Optimization Techniques

1. Multi-stage Docker builds:

```
# Build stage
FROM maven:3.9.6-eclipse-temurin-17 as build
WORKDIR /app
COPY pom.xml .
RUN mvn dependency:go-offline -B
COPY src ./src
RUN mvn package -DskipTests

# Runtime stage
FROM eclipse-temurin:17-jre
COPY --from=build /app/target/product-service.jar product-service.jar
ENTRYPOINT ["java", "-Xms256m", "-Xmx512m", "-jar", "product-service.jar"]
```

2. JVM Tuning:

- Memory allocation settings
- Garbage collection optimization
- Container-aware memory limits

3. Health Checks:

- PostgreSQL readiness check
- Kafka connectivity check
- Application health probes

Cloud Deployment Considerations

AWS Deployment

- ECS Fargate for containerized deployment
- RDS PostgreSQL for database
- MSK for managed Kafka
- Secrets Manager for credentials
- CloudWatch for monitoring and logging

Azure Deployment

- AKS for Kubernetes orchestration
- Azure Database for PostgreSQL
- Event Hubs for Kafka API compatibility
- Key Vault for secrets management
- Azure Monitor for observability

Google Cloud Deployment

- GKE for Kubernetes orchestration
- Cloud SQL for PostgreSQL
- Pub/Sub for event streaming
- Secret Manager for sensitive data
- Cloud Monitoring for observability

Scaling Considerations

1. Horizontal Scaling:

- Stateless design allows multiple instances
- Load balancing with sticky sessions not required
- Database connection pooling configured

2. Vertical Scaling:

- JVM memory allocation configuration
- Tuned thread pool sizes
- Optimized database query performance

3. Database Scaling:

- Read replicas for query distribution
- Connection pooling configuration
- Index optimization for common queries