

- Twelve-Factor App Methodology Implementation
  - 1. Codebase
    - One codebase tracked in version control, many deploys
  - 2. Dependencies
    - Explicitly declare and isolate dependencies
  - 3. Config
    - Store config in the environment
  - 4. Backing Services
    - Treat backing services as attached resources
  - 5. Build, Release, Run
    - Strictly separate build and run stages
  - 6. Processes
    - Execute the app as one or more stateless processes
  - 7. Port Binding
    - Export services via port binding
  - 8. Concurrency
    - Scale out via the process model
  - 9. Disposability
    - Maximize robustness with fast startup and graceful shutdown
  - 10. Dev/Prod Parity
    - Keep development, staging, and production as similar as possible
  - 11. Logs
    - Treat logs as event streams
  - 12. Admin Processes
    - Run admin/management tasks as one-off processes
  - Implementation Tools and Technologies

# Twelve-Factor App Methodology Implementation

---

This document details how each of the twelve factors is specifically implemented in our Healthcare Information Exchange SOA system.

## 1. Codebase

---

# One codebase tracked in version control, many deploys

- Each service (Patient, Provider, Appointment, Authentication, Audit, Analytics) has its own dedicated Git repository
- Main branches: `main`, `develop`, `feature/*`, `release/*`, `hotfix/*`
- Consistent deployment to dev, staging, and production environments from the same codebase
- Git hooks to enforce code quality and standards

## 2. Dependencies

---

### Explicitly declare and isolate dependencies

- Maven/Gradle dependency management with explicit versioning
- No reliance on system-wide packages
- All dependencies declared in `pom.xml` or `build.gradle` files
- Docker containers to isolate runtime environment
- Dependency vulnerability scanning in CI pipeline
- Separate `dependencies` and `devDependencies` clearly

## 3. Config

---

### Store config in the environment

- Environment-specific configuration stored in environment variables
- Sensitive values (passwords, API keys) stored in HashiCorp Vault
- Configuration files in YAML format with environment variable interpolation
- No hardcoded configuration values in code
- Config validation on application startup
- Centralized configuration server (Spring Cloud Config) with Git backend

## 4. Backing Services

---

# Treat backing services as attached resources

- All external services (databases, caches, message brokers) accessed via URLs
- Service connection details stored in environment variables
- Ability to swap backing services without code changes
- Health checks for all backing services
- Circuit breakers (with Resilience4j) for graceful handling of backing service failures
- Consistent database access via JPA regardless of database vendor

## 5. Build, Release, Run

---

### Strictly separate build and run stages

- CI/CD pipeline with distinct stages:
  - Build: Compile code, run tests, create artifacts
  - Release: Combine artifacts with environment config
  - Run: Execute application with immutable release
- Docker images built once, promoted through environments
- Artifact versioning with semantic versioning
- Immutable releases, no changes to deployed code
- Blue/green deployments for zero-downtime releases

## 6. Processes

---

### Execute the app as one or more stateless processes

- Stateless services that store no local state between requests
- Session data stored in Redis, not in local memory
- Persistent data always written to backing services
- No sticky sessions, enabling easy horizontal scaling
- Graceful handling of process failures
- Designed for concurrent execution of multiple instances

# 7. Port Binding

---

## Export services via port binding

- Each service exports HTTP API on configured port
- Self-contained services with embedded servers (Spring Boot)
- No reliance on external application servers
- Dynamic port allocation in container environments
- Health and metrics endpoints on dedicated ports
- API gateway for unified external access

# 8. Concurrency

---

## Scale out via the process model

- Horizontal scaling through multiple service instances
- Workload distributed through load balancing
- Kubernetes Horizontal Pod Autoscaler for automatic scaling
- Share-nothing architecture allowing concurrent instances
- Thread pools and async processing for optimized resource usage
- Event-driven architecture for parallel processing

# 9. Disposability

---

## Maximize robustness with fast startup and graceful shutdown

- Fast startup times (<10 seconds) for rapid scaling
- Graceful shutdown handling in-flight requests
- Kubernetes readiness and liveness probes
- Connection pooling for efficient resource utilization
- Crash-only software design principles
- Idempotent operations to handle retries safely

## 10. Dev/Prod Parity

---

### Keep development, staging, and production as similar as possible

- Containerized environments for consistent runtime
- Same backing services in all environments (with different scales)
- Infrastructure as Code for consistent environment provisioning
- Continuous deployment to ensure minimal drift between environments
- Development database populated with anonymized production data
- Consistent monitoring across all environments

## 11. Logs

---

### Treat logs as event streams

- Logs written to stdout/stderr
- Centralized log collection with ELK stack
- Structured logging in JSON format
- Consistent correlation IDs across service boundaries
- Log levels configurable without code changes
- Separation of application logs from audit logs

## 12. Admin Processes

---

### Run admin/management tasks as one-off processes

- Database migrations as separate processes (Flyway/Liquibase)
- Scheduled tasks using Kubernetes CronJobs
- Admin functionality exposed through secure API endpoints
- Maintenance scripts checked into version control
- Runbook documentation for operational procedures

- CLI tools for ad-hoc administrative tasks

# Implementation Tools and Technologies

---

Factor	Implementation Technologies
Codebase	GitHub, Git Flow
Dependencies	Maven/Gradle, Docker
Config	Spring Cloud Config, HashiCorp Vault
Backing Services	PostgreSQL, MongoDB, Redis, Kafka
Build, Release, Run	Jenkins/GitHub Actions, JFrog Artifactory
Processes	Spring Boot, Kubernetes
Port Binding	Spring Boot embedded servers
Concurrency	Kubernetes HPA, async processing
Disposability	Graceful shutdown hooks, readiness probes
Dev/Prod Parity	Docker, Terraform, Kubernetes
Logs	ELK Stack, Fluentd
Admin Processes	Flyway, Kubernetes Jobs